


Don't believe the hype? A hands-on introduction to machine-learning in Python

Open Workshops on Computer Based Systems Modelling

Johannes Schmidt, Johann Baumgartner

Institute for Sustainable Economic Development, BOKU, Vienna

7.05.2019



Contents - Workshop

- ▶ Day 1: Introduction to Machine Learning
- ▶ Day 2: Understanding backpropagation and Neural Networks in Python
- ▶ Day 3: An example of a practical application of Neural Networks for image recognition in Python and Reinforcement learning in Python
- ▶ Day 4: Bring your own data!

Literature

- ▶ Introduction to Machine Learning by Andrew Ng (Coursera.org)
- ▶ Deep Learning with Python by Francois Chollet

Contents - Today

- ▶ The basic concept of machine learning
- ▶ Some examples of machine learning
- ▶ Supervised, unsupervised, reinforcement learning: practical exercises
- ▶ An introduction to neural networks

Simple and hard tasks


Traditionally simple tasks for a computer (but hard for humans)

- ▶ $\log \sqrt{x^3 + y * z^7}, x = 329$
- ▶ Search for all occurrences of "European Union" in a (superlong) text
- ▶ Generate a 3D scene from a mathematical description, using raytracing
- ▶ Compile a computer program

Traditionally hard tasks for a computer (but simple for humans)

- ▶ Who is on that photo?
- ▶ Is this a cat or a dog?
- ▶ Generate a realistic image of a person
- ▶ Generate human readable text
- ▶ Translate one language to another one

The basic concept of machine learning



Why is this cool or frightening?


- ▶ Image recognition with captionbot.ai
- ▶ Computer plays Atari computer games
- ▶ The moment alphaGo wins against Lee Sedol
- ▶ This person does not exist
- ▶ Minority report
- ▶ Supervising Oktoberfest waiters
- ▶ Write a (boring) book
- ▶ Amazon automatically tracks and fires warehouse workers
- ▶ AI Lie detector for border control
- ▶ Twitter content filter going wrong
- ▶ Arbeitsmarktservice predicts probability of finding a job

A typology of machine learning


- ▶ Supervised learning: Given inputs and outputs, find the rules that link the two
- ▶ Unsupervised learning: Find structure in data
- ▶ Reinforcement learning: learning by doing.

Supervised learning - Exercise (I)

Positives



Negatives



Positive or Negative?




Supervised learning needs labeled datasets

- ▶ Classification. Input features: Photos of people. Labels: smiling or not.
- ▶ Classification. Input features: Photos of waiters carrying plates and glasses. Labels: number of plates and glasses.
- ▶ Classification. Input features: German text. Labels: English text.
- ▶ Regression. Input features: wind speeds. Labels: measured wind power generation.
- ▶ Therefore: Inmates in Finland are training AI as part of prison labor

Unsupervised learning - Exercise

Which data points belong to each other?



Unsupervised learning - Find structure in data

- ▶ Clustering (which data belongs together?)
- ▶ Anomaly detection (which data is somehow strange?)
- ▶ Generative adversarial networks (Generate photos, sounds, and text, also involves supervised learning)

Reinforcement learning - Exercise

See Netlogo

Reinforcement learning - Learning by doing

- ▶ Propose new videos on video platform to user
- ▶ Play Computer Games or Go
- ▶ Simulation of heuristic optimization of agents (e.g. on markets)

What is it?

- ▶ Google translate
- ▶ Amazon product suggestion
- ▶ Face recognition
- ▶ Facebook face recognition
- ▶ Bidding for advertisements on Taobao

An introduction to artificial neural networks

- ▶ Why neural networks?
- ▶ Basic concepts and terminology
- ▶ Backpropagation algorithm and its computational complexity
- ▶ Under- and overfitting of neural networks
- ▶ Practical aspects of machine learning

Why artificial neural networks?

- ▶ Current state of the art in pattern recognition, and language processing
- ▶ Used in supervised learning, unsupervised learning, and reinforcement learning
- ▶ Understanding neural networks will help you to understand other algorithms too.

Are they new?

- ▶ Well... not really
- ▶ 1940ies first work
- ▶ Short lived hypes in 60ies and 80ies
- ▶ 1997: US Mail uses ANNs for recognizing handwritten addresses.
- ▶ New hype starting around 2010

What are they used for?

- ▶ Classification
- ▶ Regression
- ▶ Text and Image generation
- ▶ Anything, that needs the approximation of a function.

Wait... regression? Classification?

You've heard it before, but no ANNs were involved? Well, ordinary least squares regression is around since the 19th century.

Ok, so what's the difference?

- ▶ Ordinary Least Squares regression needs assumptions on functional relationships. (blackboard example)
- ▶ ANNs are *Universal Function Approximators* for linear, bounded functions. For a nice visual proof of the theorem, see [here](#).
- ▶ ANNs **do not** allow deriving impact of input parameters on output (such as statistical significance)
- ▶ ANNs are therefore useful for prediction, but not for statistical inference
- ▶ Machine learning uses way cooler vocabulary.

Universal Function Approximator?




Possible Input Output relations

- ▶ German Language -> English Language
- ▶ Photos of cats and dogs -> Classification "Cat", "Dog"
- ▶ $x \rightarrow x^2$
- ▶ Weather data -> Wind power generation

Elephants

Wait... but what about "With four parameters I can fit an elephant, and with five I can make him wiggle his trunk." (John von Neumann)



Fit the elephant in python with four complex numbers

So true: the art of ANNs lies in achieving reasonable fits without **overfitting**.

Basic terminology

- ▶ Input layer (features)
- ▶ Hidden layer
- ▶ Output layer (labels)
- ▶ Nodes
- ▶ Connections
- ▶ Weights
- ▶ Bias node
- ▶ Activation function
- ▶ Error measure

Backpropagation algorithm

- ▶ Weights are adapted, so that the prediction and the observation error is minimized.
- ▶ Non-convex optimization due to non-convex activation functions: heuristic approach
- ▶ Training neural networks is NP-Complete in its most general form.

Terminology data sets and training

- ▶ Sample
- ▶ Training data
- ▶ Validation data
- ▶ Test data
- ▶ Epoch
- ▶ Batch, batch size
- ▶ Online training, micro-batch training, batch training

Underfitting

- ▶ Both training and validation data do not fit well.
- ▶ Problem: not enough data to fit to your network
- ▶ Solution: Smaller network. Or more data.

Overfitting

- ▶ Training data fits excellent, validation or test data does not.
- ▶ Problem: Network is adapted perfectly to your training data, but does not generalize.
- ▶ Solution: larger network. Add leaky ReLu layers.

Practical aspects

- ▶ Always assess variance-bias trade-off to understand under- and overfitting
- ▶ Get a GPU!
- ▶ Play around (get a GPU!)
- ▶ Choose your problems wisely. (If OLS regression works, use it, it gives much more insights.)
- ▶ Use feature engineering. With and without neural networks.

Next workshops

What we do

- ▶ Play around in netlogo to better understand backpropagation
- ▶ How to predict windpower from wind speeds
- ▶ Practical exercises in python

What to bring

- ▶ A laptop
- ▶ Python installed with the following packages:
- ▶ Netlogo installed

Thank you!

For slides and source-code, check

<https://github.com/joph/Machine-Learning-Workshop>

mail: johannes.schmidt@boku.ac.at

