

Don't believe the hype? A hands-on introduction to machine-learning in Python - Part III

Open Workshops on Computer Based Systems Modelling

Johannes Schmidt, Johann Baumgartner

Institute for Sustainable Economic Development, BOKU, Vienna

7.05.2019



Contents

- ▶ A word on the blackbox
- ▶ Pattern recognition with neural networks: some work in progress
- ▶ Q-Learning

A word on the blackbox

- ▶ **Partial dependence plot:** Average effect of one variable on output, over complete parameter space. Shows functional relationship between input and output. Assumption: inputs are all uncorrelated.
- ▶ Other methods are available (e.g. Accumulated Local Effects plot, Individual Conditional Expectation)
- ▶ **LSTM-Vis:** visualize hidden states of LSTMs. Interesting to understand, what the hidden states may represent. See also [this](#) for an example of what hidden states represent in hydrology.

Pattern recognition with neural networks: work in progress

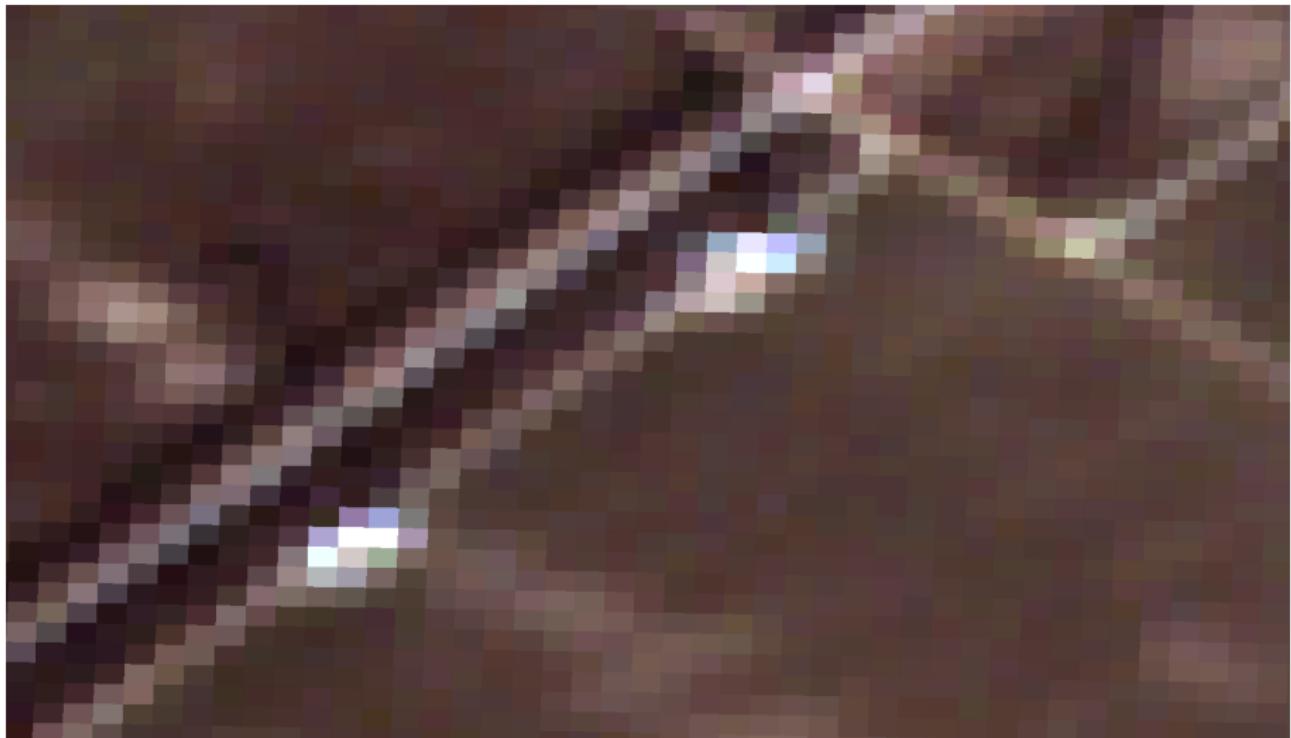
- ▶ Validating the Global Power Plant Database with satellite fotos...
- ▶ Work in progress. Results not extremely promising.
Recommendations very much welcome.

Global Power Plant Database: Need for Validation



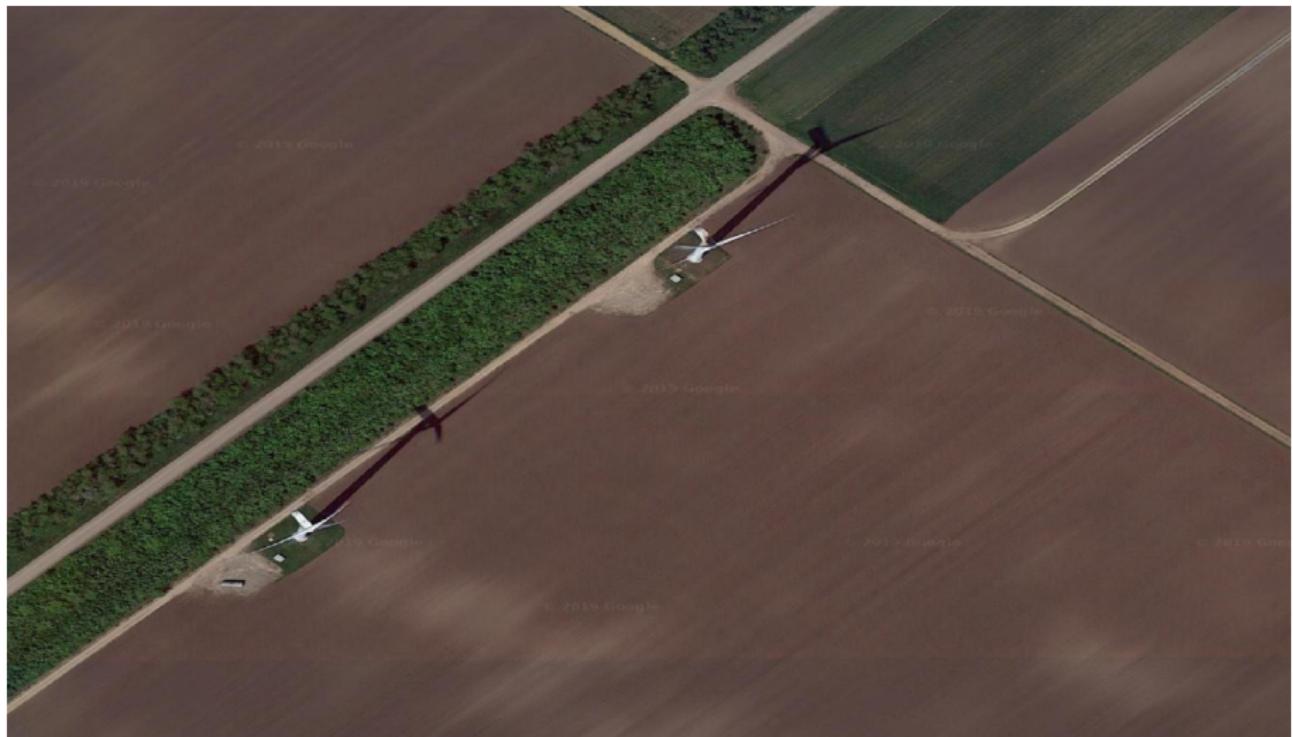
Satellite data for validation

Sentinel-2 (10 meter resolution), free

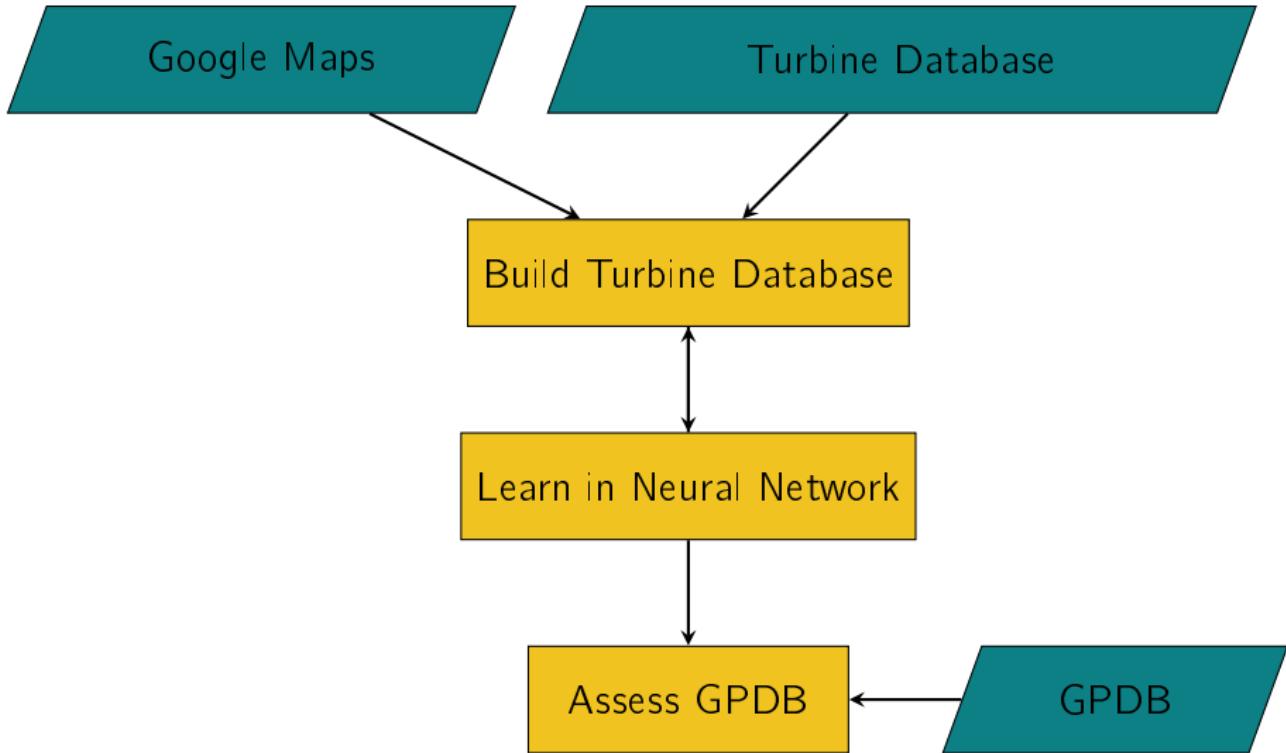


Satellite data for validation

WorldView-4, Quickbird (0.31-0.65 meter resolution) used by e.g. Google Maps



Approach



Software

- ▶ Downloading of data from google maps: R (RGdal, tidyverse, raster)
- ▶ Machine Learning: Python (keras, scikit-image, gdal)
- ▶ Mixing R and Python: not a brilliant idea. Just lazy.

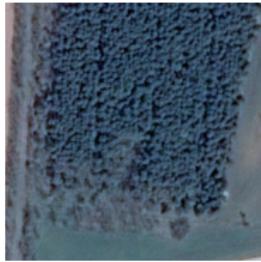
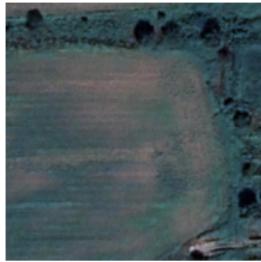
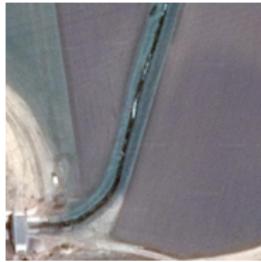
Create Sample Database

Manual quality control

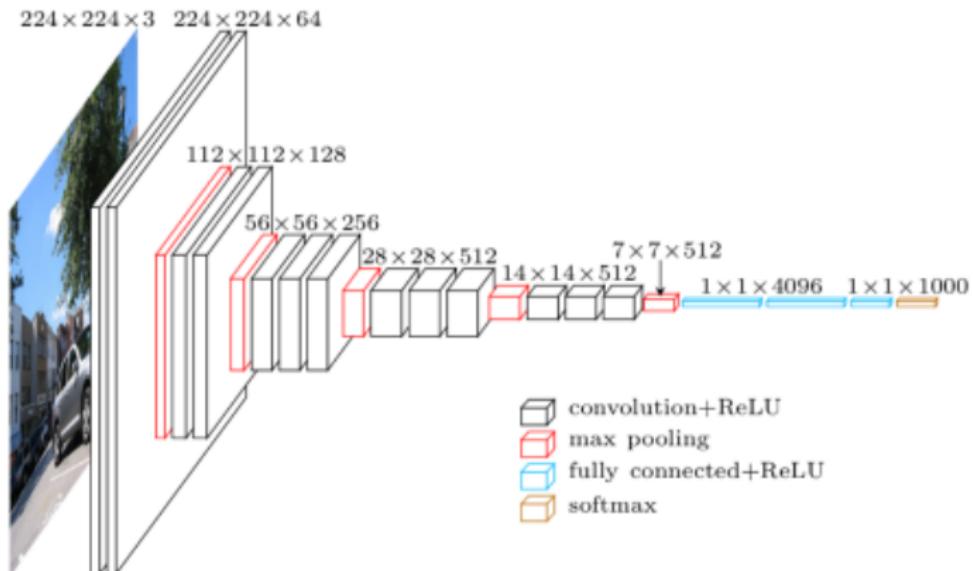
Positives



Negatives



Learn with Convolutional layers



Layered feature learning

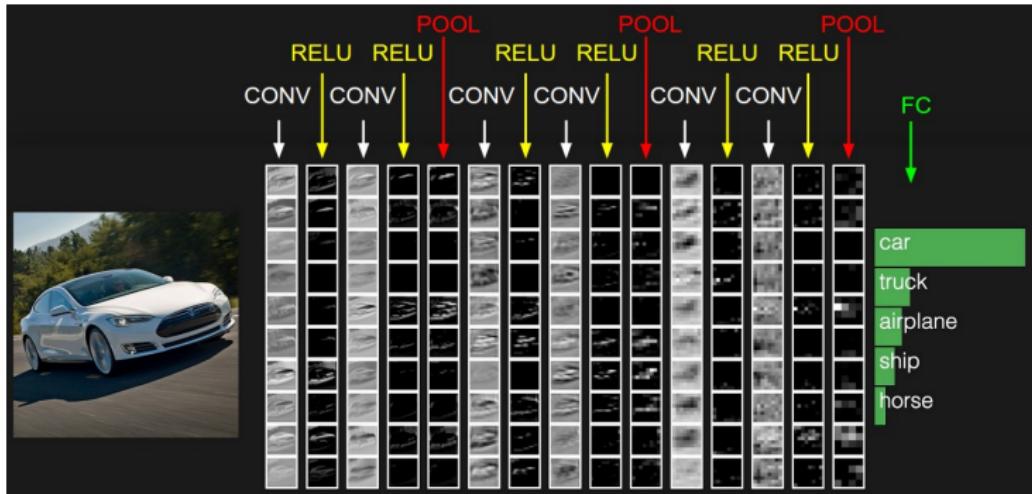


Figure: Taken from [here](#)

Terminology

- ▶ Convolution operation
- ▶ Filter
- ▶ Depth
- ▶ Stride
- ▶ Zero padding (wide and narrow convolution)

Pooling

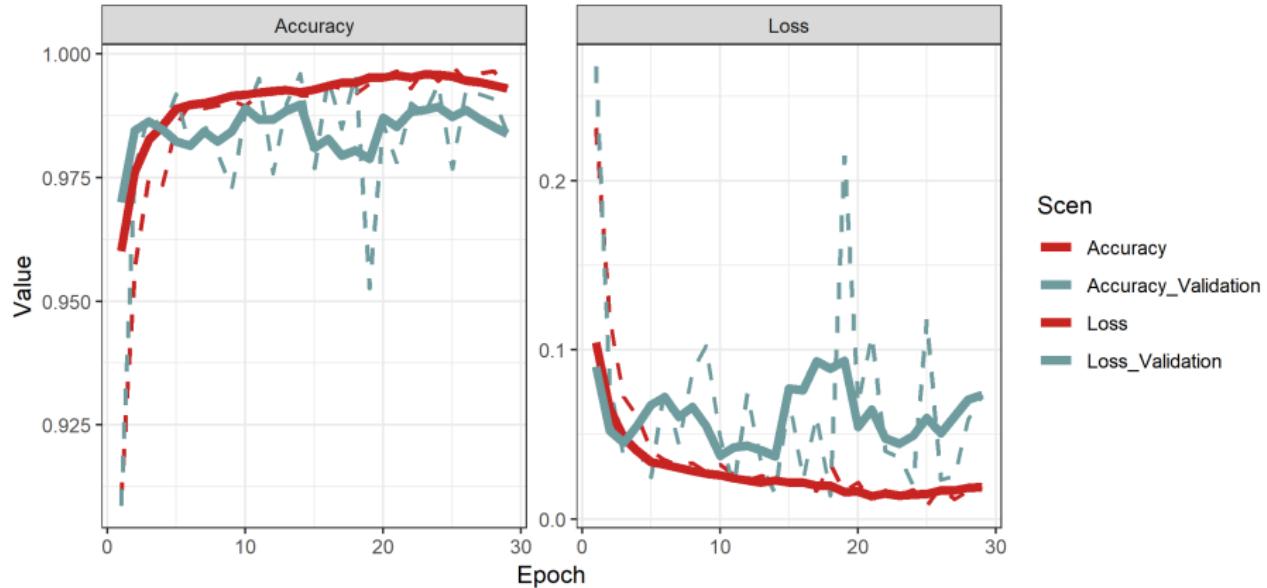
- ▶ Makes input representations (feature dimension) smaller and more manageable
- ▶ Reduces the number of parameters and computations in the network controlling overfitting
- ▶ Makes the network invariant to small transformations, distortions and translations in the input image
- ▶ Helps us arriving at an almost scale invariant representation of our image (the exact term is “equivariant”)
- ▶ taken from [here](#)

Some hints

- ▶ Compile high quality training set
- ▶ Do not overfit!
- ▶ Use pretrained networks to speed up learning process. Keras models can be found [here](#)
- ▶ Unfreeze last layers to learn your specific task, fixing the lower layers which extract basic features

Results Learning

Quality Measures



$$\text{Accuracy} = \frac{\text{Number of Correct Predictions}}{\text{Number of Total Predictions}}$$

$$\text{Loss} = -1 \frac{1}{N} \sum_{i=1}^N N y_i \log(p(y_i)) + (1 - y_i) \log(1 - p(y_i))$$

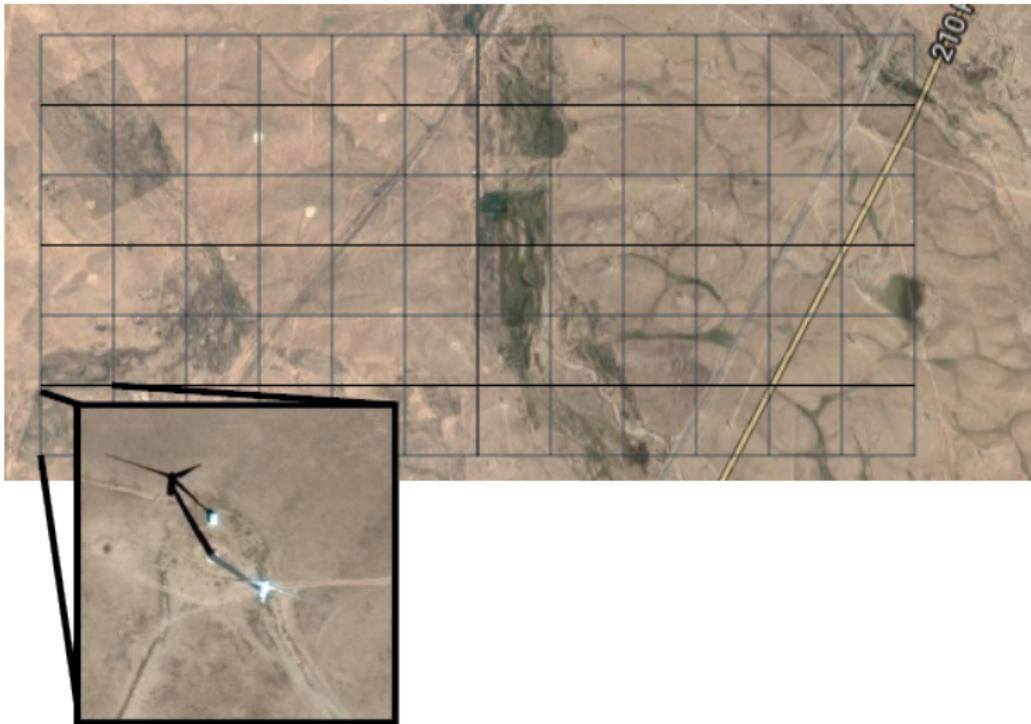
Results Learning

Class Activation Map

see notebook

Searching wind turbines

Create grid and check image by image



Application

French Wind Parks

France: 10 parks assessed.

| Turbines correctly identified | Non-Turbines correctly identified | Turbines wrongly identified | Non-Turbines wrongly identified |
|-------------------------------|-----------------------------------|-----------------------------|---------------------------------|
| 6 | 3429 | 5 | 8 |

Discussion

- ▶ Beware: Shaky legal conditions. Also, date of satellite photos unknown.
- ▶ Runtime prohibitive on desktop computers. For a full global check would need cluster computing with high bandwidth (but see shaky legal conditions...)
- ▶ Improve classification: need for more negatives which are similar to positives? How could we derive those? (Currently random sampling of negatives, therefore tiles do not contain a lot of cities or roads...)

Conclusions

- ▶ Full validation of GPDB due to missing temporal information not fully possible.
- ▶ Current best commercial satellite data allows identification of single turbines with good accuracy by using machine learning approaches
- ▶ Research limited by Data availability. Public domain data is low resolution (i.e. Sentinel-2) or small spatial domain (free ortho-photos, like basemmap.at).

And now for something completely different...

Contents

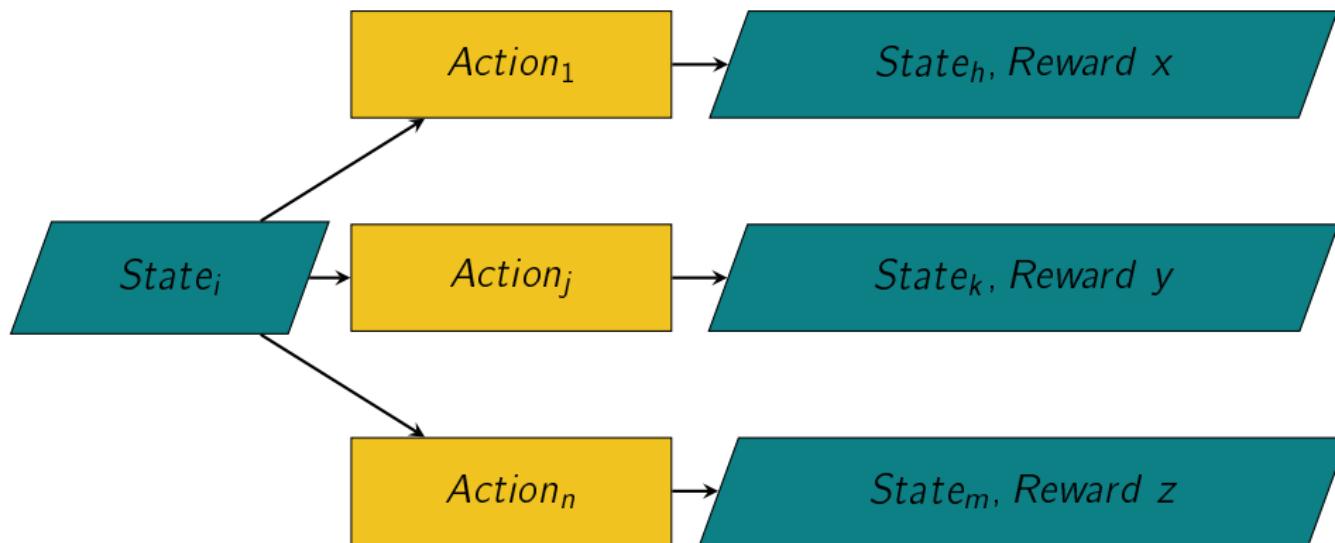
- ▶ The Q-learning algorithm
- ▶ A very simple example calculated by hand
- ▶ A simple example in python

Some examples

- ▶ Pan flipping robot
- ▶ The moment alphaGo wins against Lee Sedol
- ▶ Automatic recommendation systems

The Q-learning algorithm in a nutshell

Which action to pick?



The Q-Table

The Q-Table is updated in a learning by doing approach

Q-Updating rule introduced 1989 by Watson. Updates Q-Table according to current rewards and expected future rewards

$$q^{new}(s_t, a_t) = (1 - \alpha)q(s_t, a_t) + \alpha(r_t + \gamma \max_a q(s_{t+1}, a))$$

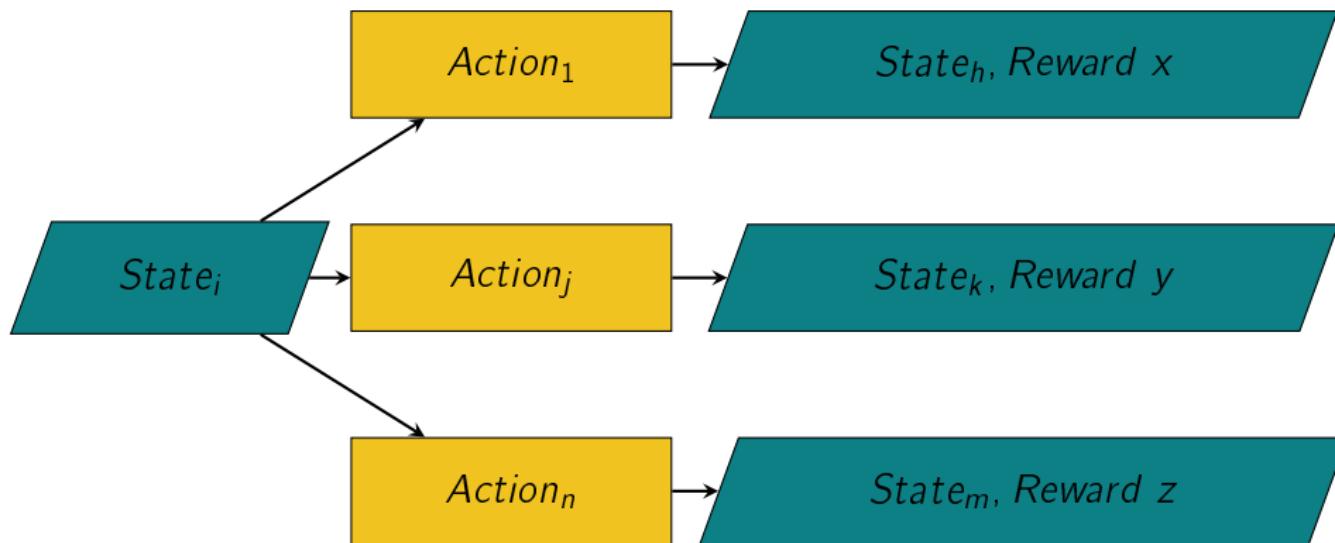
α is learning rate, γ is discount rate

r_t is reward, s_t is state, a_t is action

| | | Actions | | | |
|--------|--|----------------|----------------|----------------|----------------|
| | | q_{s_1, a_1} | q_{s_1, a_2} | q_{s_1, a_3} | q_{s_1, a_4} |
| States | | q_{s_2, a_1} | q_{s_2, a_2} | q_{s_2, a_3} | q_{s_2, a_4} |
| | | q_{s_3, a_1} | q_{s_3, a_2} | q_{s_3, a_3} | q_{s_3, a_4} |
| | | q_{s_4, a_1} | q_{s_4, a_2} | q_{s_4, a_3} | q_{s_4, a_4} |

The Q-learning algorithm in a nutshell

Which action to pick?



The Q-Table

The Q-Table is updated in a learning by doing approach

Q-Updating rule introduced 1989 by Watson. Updates Q-Table according to current rewards and expected future rewards

$$q^{new}(s_t, a_t) = (1 - \alpha)q(s_t, a_t) + \alpha(r_t + \gamma \max_a q(s_{t+1}, a))$$

α is learning rate, γ is discount rate

r_t is reward, s_t is state, a_t is action

| | | Actions | | | |
|--------|--|----------------|----------------|----------------|----------------|
| | | q_{s_1, a_1} | q_{s_1, a_2} | q_{s_1, a_3} | q_{s_1, a_4} |
| States | | q_{s_2, a_1} | q_{s_2, a_2} | q_{s_2, a_3} | q_{s_2, a_4} |
| | | q_{s_3, a_1} | q_{s_3, a_2} | q_{s_3, a_3} | q_{s_3, a_4} |
| | | q_{s_4, a_1} | q_{s_4, a_2} | q_{s_4, a_3} | q_{s_4, a_4} |

Choosing an action

- ▶ If random number $< \epsilon$, take a random action (exploration)
- ▶ Otherwise, take the action with the highest q -value (exploitation)

Examples

- ▶ A simple game on the blackboard
- ▶ Navigate a maze with q-learning

Thank you!

For slides and source-code, check

<https://github.com/joph/Machine-Learning-Workshop>

mail: johannes.schmidt@boku.ac.at

